



Jev + CLAUDE

the setup guide

A ROBONUGGETS GUIDE

The promise: pair Jev with Claude and get 10x the power at 100x less cost. Jev makes the fast calls for a fraction of a cent. Claude keeps the thinking and the writing.

What Jev is, in 4 lines

1	Not an LLM	It never writes text. You send it a state (any text or data) plus typed questions, and it sends back decisions.
2	3 answer shapes	choice picks one of your options. score places it on your scale. noul gives the chance a statement is true.
3	Output is free	You pay \$0.042 per 1M input tokens and nothing at all for the answers.
4	About 0.3s	It answers every question in one parallel pass, in 70 to 500 ms.

Jev is the first "System One" model from TypeSafe AI, launched 15 September 2026.

README

What this is. The setup guide for the "Jev + Claude" lesson. It shows how to wire Jev, a fast decision model, into Claude Code, then three levels of things to do with it.

How to set it up. Work through the pages in order. This page gets your key. Page 3 is the one prompt that wires Jev in. Pages 4 and 5 are the builds on top: the routers, then sorting any pile of text. Page 6 shows where this goes next. Anything in a bordered box is copy-paste ready.

Two ways to get this. Build it yourself with the prompts in this guide, free. Or install the pre-built kit from the video: the router hook, the four helper agents and the /jev command. That comes with the RoboNuggets community at skool.com/robonuggets.

01 Get access

Jev is open to everyone. The simplest way in is OpenRouter: one account and one key that reaches many AI models, Jev included.

1. **Make an OpenRouter account** at openrouter.ai.
2. **Add a little credit.** Jev is a paid model, but a tiny one. Our whole 100-email test on page 5 cost **\$0.0032**.
3. **Create an API key** at openrouter.ai/keys and copy it. An API key is a password that lets Claude call Jev for you. Treat it like one.
4. **Keep it handy for page 3.** The prompt asks for it and keeps it somewhere safe on your computer. You set nothing up by hand.

MODEL ID

typesafe/jev-1.13

The name Claude uses when it calls Jev through OpenRouter.

LIMITS

64K tokens per request, with 32K of that for your data. Text only. English works best.

What it costs

MODEL	INPUT, PER 1M TOKENS	OUTPUT, PER 1M TOKENS	INPUT PRICE VS JEV
Jev	\$0.042	free	1x
Claude Haiku 4.5	\$1	\$5	24x
Claude Fable 5.1	\$10	\$50	238x

List prices in USD, checked 21 September 2026. Prices move, so check the provider's own page before you rely on them.

Other ways in, if you already use them: Vercel AI Gateway ([typesafe-ai/jev](https://vercel.com/docs/ai)), Cloudflare Workers AI ([typesafe/jev](https://workers.cloudflare.com/)), or TypeSafe direct (keys at console.typesafe.ai). This guide uses OpenRouter throughout.

What one Jev call looks like

YOU SEND A STATE

One email

"Hi, we run an online skincare store and want an AI agent to answer our order emails. Budget is signed off and we want to start next month. Can we book a call this week?"



YOU ASK

3 typed questions

choice: which team should reply?
score: how good a lead is it?
noul: does it need a personal reply today?



JEV ANSWERS

3 decisions

sales, confidence 1
2.99 on a 0 to 3 scale (hot lead),
confidence 0.99
0.78 chance it needs a reply today

A real call through OpenRouter, recorded 21 September 2026: 0.26 s (the middle of 3 runs) for \$0.00002. The prompt on page 3 runs this same test for you.

02 Wire it into your agent

Open your AI agent in an empty folder or in your existing workspace and paste this whole prompt. We use Claude Code in the video, and the prompt works in any agent that can run things on your computer. Your agent asks for your key, makes one real test call so you can watch Jev answer, then saves it all as a reusable skill. Jev launched after every model's training data, which is why the prompt points at the public docs instead.

COPY-PASTE · THE MAIN PROMPT

i want to add a new kind of AI model called Jev to this workspace. Jev doesn't write anything. you hand it some text plus a few questions, and it answers them almost instantly, for almost nothing. it's made for sorting and deciding.

- Jev is brand new, so you won't know it yet. before anything else, read how it works:

<https://docs.typesafe.ai/llms.txt>

- i'll use it through OpenRouter, where it's called typesafe/jev-1.13. how to call it there:

<https://openrouter.ai/docs/api/api-reference/alphadecisions/submit-a-decisions-questions-and-answers-request>. it's a brand new route, so follow those docs if anything has moved

- ask me for my OpenRouter key (i get one at <https://openrouter.ai/keys>). keep it somewhere safe on my computer, never inside a file i might share, and never show it back to me

- Jev can only answer in 3 shapes: pick one option from a list, give a score on a scale, or say how likely something is true. most answers also come back with how sure it is

- do one real test so i can watch it work. make up a short sales email and ask Jev 3 things: how strong a lead is this, what kind of email is it, and does it need a personal reply. show me its answers, how long it took and what it cost. if the test fails, show me the exact error


- then save what you learned as a reusable skill (or whatever your setup calls a saved instruction), so next time i can say "use Jev to sort these" about any pile of text

the rule from here on: Jev decides, you write. when Jev isn't sure about something, you make the call yourself. anything you send to Jev leaves my computer, so ask me before sending anything private.


When your agent finishes, try it on any pile of text: "use Jev to sort these".

03 Level 1 – route every message

Most messages do not need the biggest model. Jev reads each one and names the smallest Claude model that can do it.



The honest constraint. Claude Code has no built-in per-message model switch, and a hook cannot change the model. A hook can only add a note. So the router is a note plus four helper agents, each pinned to one model by the model: line in its file. The prompt asks your agent to be honest about this, which stops it inventing a setting that does not exist.



It never blocks a message, and it goes silent if Jev is slow or the key is missing. While it is on, your message text goes to TypeSafe through OpenRouter. It works in the Claude Code VS Code extension, the terminal, and the Code tab of the Claude desktop app.

COPY-PASTE · THE MODEL ROUTER PROMPT

build me a model router that uses Jev. the idea: Jev sizes up every message i send, and small jobs go to a smaller, cheaper AI model instead of the biggest one. Jev is already set up in this workspace.

- for each message, ask Jev one question: what's the smallest model that can do this job well? give it 4 sizes to pick from: tiny (a lookup, a rename, a one-line answer), everyday (a normal email, post or short document), large (a multi-step build, research, a full report) and hardest (strategy, or anything where a wrong call is expensive)
- match each size to a model i actually have, smallest to biggest, and make one helper agent per size that runs on that model. if i have fewer models than sizes, use fewer sizes. each helper ends its reply with one line saying which model did the work
- when Jev is sure, hand the job to the matching helper and pass the result back to me. when Jev is less than 60% sure, or my message is a short reply that only makes sense inside our conversation, handle it yourself
- it must never slow me down or block a message. if Jev is slow, or anything at all goes wrong, carry on as if the router wasn't there
- give me a simple on and off switch, plus a status check that shows how many messages went to each size and what Jev has cost so far. keep it OFF until i turn it on
- be honest with me: if your setup can't switch models per message, say so, and build the closest thing that works
- test it for real with 8 sample messages, from tiny to hardest, plus 2 short replies like "yes do that but make it shorter". show me a table of what Jev picked and how sure it was

leave it OFF when you're done and tell me the exact words to turn it on. remind me that while it's on, my messages pass through OpenRouter to the company that makes Jev, so it stays off for private work.

The skill router, same idea. Jev reads the name and first line of every skill and returns the one to load, with its confidence. Under 0.6, Claude picks. In our recorded run Jev chose the right skill 12 times out of 14, from a list of 145. TypeSafe's own cookbook test reports wrong-skill loads falling from 17% to 7.3% (182 skills, agent on Haiku 4.5). That is their number, not an independent one.

04 Level 2 - sort anything

BEFORE

A person opens every email, ticket or invoice, one at a time. Or you pay a big chat model to read the whole pile.

AFTER

Jev sorts the whole pile in seconds. Claude only writes the replies that matter.

What we measured: the same 100 emails, three models

MODEL	TIME	COST	LEAD SCORE EXACTLY RIGHT	HOT LEADS FOUND
Jev	3.6 s	\$0.0032	89 / 100	15 / 15
Claude Haiku 4.5 (thinking off)	11.1 s	\$0.061	84 / 100	15 / 15
Claude Fable 5.1 (effort low)	38.7 s	\$0.76	97 / 100	15 / 15

Fable was the most accurate, Jev beat Haiku, and all three found all 15 hot leads.

100 made-up emails for a fictional agency. Same questions, 10 at a time, real calls through OpenRouter, recorded 21 September 2026. Cost is what OpenRouter billed. Also from the video: 27 questions about one support ticket, answered in one pass in about 0.25 s for \$0.00008.

COPY-PASTE · SORT MY INBOX EXPORT BY LEAD QUALITY

use Jev to sort my inbox by lead quality. the emails are in [my export file or folder]. ask Jev 3 things about every email:

- how strong a lead is this? 4 levels: not a lead at all (spam, vendors, newsletters, job seekers, support, existing clients), cold (vague interest, tiny budget, poor fit), warm (a real need that fits, but no budget or timeline yet), hot (a clear need that fits, plus at least two of: a stated budget, a stated timeline, a decision maker writing)
- what kind of email is it? new lead, existing client, vendor pitch, spam, job seeker, newsletter or support
- does it need a personal reply from my team? yes or no
- tell Jev in one line what my business does, send the emails in small batches so it stays fast, and give me one spreadsheet with the hot leads at the top

anything Jev is less than 60% sure about goes in a "check these" pile for you to judge. then draft replies for the hot leads only. writing is your job, not Jev's.

COPY-PASTE · TRIAGE SUPPORT TICKETS BY URGENCY AND TEAM

use Jev to sort my support tickets by urgency and by team. the tickets are in [my export file, or i'll paste them]. ask Jev 3 things about every ticket:

- how soon does it need an answer? today, this week, or no rush
- which team should take it? technical (bugs, outages, errors), billing (charges, refunds, invoices), sales (upgrades, pricing) or success (onboarding, training)
- how likely is this customer to leave? 4 levels: no sign of leaving, mild frustration, openly looking at other options, has set a deadline to leave
- if i have the customer's plan and how long they've been with us, send that along with the ticket. send tickets in small batches so it stays fast, and give me one table sorted by urgency, then by how likely they are to leave

anything Jev is less than 60% sure about gets flagged for a person to read. then write a first reply for the "today" tickets only. writing is your job, not Jev's.

COPY-PASTE · CHECK SUPPLIER INVOICES FOR FRAUD SIGNS

use Jev to check my supplier invoices for signs of fraud before anyone pays them. the invoices are in [my folder]. ask Jev 3 things about every invoice:

- how risky does it look? 4 levels: looks normal, one small oddity, several warning signs, strong signs of fraud
- does it combine changed bank details with pressure to pay fast? yes or no
- what should we do with it? pay as normal, hold it and phone the supplier on a number we already know, or reject it
- Jev reads text only, so turn each invoice into plain text first. send it together with what we know about that supplier: their usual bank details, their usual amounts and the date of their last invoice
- Jev is weak at maths and at comparing dates, so you do the sums and the date checks yourself. only ask Jev the judgement calls

give me a table with the riskiest invoices first and one line on why each was flagged. nothing gets paid or rejected on Jev's word alone. it only decides what a person looks at first.

05 Level 3 – what this unlocks

Once a decision costs almost nothing and takes a third of a second, you can put one in places where a chat model was always too slow. One real example is already out.

UNCLUTTER, IN 5 LINES

- A free, open-source browser extension by Kitze that cleans up web pages: github.com/kitze/unclutter
- Jev judges up to 60 page elements in one pass: keep it, or is it an ad, a promo, a newsletter box, a social widget or a cookie banner?
- It only hides what Jev is at least 0.9 sure about. Anything uncertain stays on the page.
- It saves the rule per page type, so the next article on that site is cleaned with no new call.
- It takes a Vercel AI Gateway key or a TypeSafe key, not an OpenRouter key.

When NOT to use Jev

Writing

Chat

Reasoning

Counting

Maths

Dates

This is TypeSafe's own list of weak spots, plus long input full of irrelevant text. Give those jobs to Claude. The rule from page 3 holds everywhere: Jev decides, Claude writes.

06 Where this goes next

The easy part is wiring Jev in. One prompt does it. The skill that pays is spotting which decisions in a business are worth handing to a fast, cheap model, then building the system around them.

That is what we practise every week: building AI systems you can sell, and mastering Claude Code + Claude Design along the way. The pre-built Jev kit from the video lives there too, ready to install.

join us at

skool.com/robonuggets

Every link in this guide

openrouter.ai – account and credit

github.com/kitze/unclutter – Unclutter

skool.com/robonuggets – the kit and the community

openrouter.ai/keys – your API key

console.typesafe.ai – TypeSafe direct keys

Model id: `typesafe/jev-1.13`